

```
In [1]: import pandas as pd  
import matplotlib.pyplot as plt  
import numpy as np
```

```
In [2]: df=pd.read_csv("/home/rmdstic/Documents/TEA-14/Social_Network_Ads.csv")
df
```

Out[2]:

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|-----|----------|--------|-----|-----------------|-----------|
| 0 | 15624510 | Male | 19 | 19000 | 0 |
| 1 | 15810944 | Male | 35 | 20000 | 0 |
| 2 | 15668575 | Female | 26 | 43000 | 0 |
| 3 | 15603246 | Female | 27 | 57000 | 0 |
| 4 | 15804002 | Male | 19 | 76000 | 0 |
| 5 | 15728773 | Male | 27 | 58000 | 0 |
| 6 | 15598044 | Female | 27 | 84000 | 0 |
| 7 | 15694829 | Female | 32 | 150000 | 1 |
| 8 | 15600575 | Male | 25 | 33000 | 0 |
| 9 | 15727311 | Female | 35 | 65000 | 0 |
| 10 | 15570769 | Female | 26 | 80000 | 0 |
| 11 | 15606274 | Female | 26 | 52000 | 0 |
| 12 | 15746139 | Male | 20 | 86000 | 0 |
| 13 | 15704987 | Male | 32 | 18000 | 0 |
| 14 | 15628972 | Male | 18 | 82000 | 0 |
| 15 | 15697686 | Male | 29 | 80000 | 0 |
| 16 | 15733883 | Male | 47 | 25000 | 1 |
| 17 | 15617482 | Male | 45 | 26000 | 1 |
| 18 | 15704583 | Male | 46 | 28000 | 1 |
| 19 | 15621083 | Female | 48 | 29000 | 1 |
| 20 | 15649487 | Male | 45 | 22000 | 1 |
| 21 | 15736760 | Female | 47 | 49000 | 1 |
| 22 | 15714658 | Male | 48 | 41000 | 1 |
| 23 | 15599081 | Female | 45 | 22000 | 1 |
| 24 | 15705113 | Male | 46 | 23000 | 1 |
| 25 | 15631159 | Male | 47 | 20000 | 1 |
| 26 | 15792818 | Male | 49 | 28000 | 1 |
| 27 | 15633531 | Female | 47 | 30000 | 1 |
| 28 | 15744529 | Male | 29 | 43000 | 0 |
| 29 | 15669656 | Male | 31 | 18000 | 0 |
| ... | ... | ... | ... | ... | ... |
| 370 | 15611430 | Female | 60 | 46000 | 1 |
| 371 | 15774744 | Male | 60 | 83000 | 1 |
| 372 | 15629885 | Female | 39 | 73000 | 0 |

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|-----|----------|--------|-----|-----------------|-----------|
| 373 | 15708791 | Male | 59 | 130000 | 1 |
| 374 | 15793890 | Female | 37 | 80000 | 0 |
| 375 | 15646091 | Female | 46 | 32000 | 1 |
| 376 | 15596984 | Female | 46 | 74000 | 0 |
| 377 | 15800215 | Female | 42 | 53000 | 0 |
| 378 | 15577806 | Male | 41 | 87000 | 1 |
| 379 | 15749381 | Female | 58 | 23000 | 1 |
| 380 | 15683758 | Male | 42 | 64000 | 0 |
| 381 | 15670615 | Male | 48 | 33000 | 1 |
| 382 | 15715622 | Female | 44 | 139000 | 1 |
| 383 | 15707634 | Male | 49 | 28000 | 1 |
| 384 | 15806901 | Female | 57 | 33000 | 1 |
| 385 | 15775335 | Male | 56 | 60000 | 1 |
| 386 | 15724150 | Female | 49 | 39000 | 1 |
| 387 | 15627220 | Male | 39 | 71000 | 0 |
| 388 | 15672330 | Male | 47 | 34000 | 1 |
| 389 | 15668521 | Female | 48 | 35000 | 1 |
| 390 | 15807837 | Male | 48 | 33000 | 1 |
| 391 | 15592570 | Male | 47 | 23000 | 1 |
| 392 | 15748589 | Female | 45 | 45000 | 1 |
| 393 | 15635893 | Male | 60 | 42000 | 1 |
| 394 | 15757632 | Female | 39 | 59000 | 0 |
| 395 | 15691863 | Female | 46 | 41000 | 1 |
| 396 | 15706071 | Male | 51 | 23000 | 1 |
| 397 | 15654296 | Female | 50 | 20000 | 1 |
| 398 | 15755018 | Male | 36 | 33000 | 0 |
| 399 | 15594041 | Female | 49 | 36000 | 1 |

400 rows × 5 columns

```
In [3]: df.drop("Gender",axis=1,inplace=True)
df.head()
```

```
Out[3]:
```

| | User ID | Age | EstimatedSalary | Purchased |
|---|----------|-----|-----------------|-----------|
| 0 | 15624510 | 19 | 19000 | 0 |
| 1 | 15810944 | 35 | 20000 | 0 |
| 2 | 15668575 | 26 | 43000 | 0 |
| 3 | 15603246 | 27 | 57000 | 0 |
| 4 | 15804002 | 19 | 76000 | 0 |

```
In [4]: df.isnull().sum()
```

```
Out[4]: User ID      0
Age      0
EstimatedSalary  0
Purchased  0
dtype: int64
```

```
In [5]: x=df.drop("Purchased",axis=1)
x.head()
```

```
Out[5]:
```

| | User ID | Age | EstimatedSalary |
|---|----------|-----|-----------------|
| 0 | 15624510 | 19 | 19000 |
| 1 | 15810944 | 35 | 20000 |
| 2 | 15668575 | 26 | 43000 |
| 3 | 15603246 | 27 | 57000 |
| 4 | 15804002 | 19 | 76000 |

```
In [6]: y=df.Purchased
y.head()
```

```
Out[6]: 0    0
1    0
2    0
3    0
4    0
Name: Purchased, dtype: int64
```

```
In [7]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest =train_test_split(x, y, test_size =0.2,
```

```
In [8]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
xtrain = sc.fit_transform(xtrain)
xtest = sc.transform(xtest)
xtest
```

/home/rmdstic/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/data.py:645: DataConversionWarning: Data with input dtype in t64 were all converted to float64 by StandardScaler.

return self.partial_fit(X, y)

/home/rmdstic/anaconda3/lib/python3.7/site-packages/sklearn/base.py:464: DataConversionWarning: Data with input dtype int64 were all converted to float64 by StandardScaler.

return self.fit(X, **fit_params).transform(X)

/home/rmdstic/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:4: DataConversionWarning: Data with input dtype int64 were all converted to float64 by StandardScaler.

after removing the cwd from sys.path.

```
Out[8]: array([[ 4.88966473e-01, -7.98950822e-01,  4.94607583e-01],
 [ -5.50592839e-01, -2.12648508e-02, -5.77359062e-01],
 [  1.21526737e+00, -3.12897090e-01,  1.46942725e-01],
 [  6.70428177e-01, -7.98950822e-01,  2.62831011e-01],
 [ -7.54886252e-01, -3.12897090e-01, -5.77359062e-01],
 [  1.74256773e+00, -1.09058306e+00, -1.44652121e+00],
 [ -1.15026219e+00, -7.01740076e-01, -1.59138156e+00],
 [ -1.50619343e-01, -2.15686344e-01,  2.14601566e+00],
 [ -3.55608066e-01, -1.96547978e+00, -5.58617754e-02],
 [ -9.75185983e-01,  8.53631867e-01, -7.80163563e-01],
 [ -2.52081338e-02, -7.98950822e-01, -6.06331134e-01],
 [  1.54471657e+00, -9.93372315e-01, -4.32498705e-01],
 [  1.16706395e+00, -1.18475597e-01, -4.32498705e-01],
 [ -1.58913047e+00,  7.59458956e-02,  2.04886868e-01],
 [  7.79563458e-01, -1.77105829e+00,  4.65635512e-01],
 [  1.38930771e+00, -6.04529329e-01,  1.36376973e+00],
 [  9.84339331e-01, -1.18475597e-01,  2.04886868e-01],
 [  1.65607968e+00, -1.86826903e+00,  4.36663440e-01],
 [ -5.24625142e-01,  1.63131784e+00,  1.74040666e+00],
 [  2.41606414e-01, -3.12897090e-01, -1.38857706e+00],
 [ -1.15717272e+00, -3.12897090e-01, -6.64275277e-01],
 [  7.80471618e-01,  8.53631867e-01,  2.14601566e+00],
 [ -1.84944950e-01,  2.70367388e-01, -5.48386991e-01],
 [  5.49245588e-01,  8.53631867e-01,  1.01610487e+00],
 [ -1.08952899e+00, -1.47942605e+00, -1.21474464e+00],
 [  8.89138629e-01,  1.04805336e+00,  2.05909944e+00],
 [ -3.39658507e-01, -9.93372315e-01,  4.94607583e-01],
 [  9.20154463e-02, -8.96161568e-01,  2.91803083e-01],
 [ -8.53265514e-01, -1.18475597e-01, -2.29694204e-01],
 [ -1.51345521e+00, -6.04529329e-01,  4.65635512e-01],
 [ -1.47843429e+00, -1.67384754e+00,  5.23579655e-01],
 [ -7.60704152e-01, -1.18475597e-01,  2.62831011e-01],
 [  1.75035804e+00,  1.82573933e+00, -2.87638347e-01],
 [ -9.79542313e-01, -1.18475597e-01, -4.90442848e-01],
 [  3.81278572e-01, -1.38221530e+00, -3.45582490e-01],
 [  7.88049077e-01, -1.96547978e+00, -5.19414919e-01],
 [ -9.08226882e-02, -1.57663679e+00,  3.20775154e-01],
 [ -1.55447849e+00, -4.10107836e-01, -7.80163563e-01],
 [ -2.53071135e-01, -7.01740076e-01, -1.04091221e+00],
 [ -1.37870698e+00,  1.04805336e+00, -9.82968063e-01],
 [  1.06549193e+00, -1.09058306e+00,  5.23579655e-01],
```

```
[ -2.39391976e-01,  2.70367388e-01, -5.19414919e-01],
[ -1.32190441e+00, -1.09058306e+00,  4.07691369e-01],
[  1.69914633e+00, -3.12897090e-01, -1.44652121e+00],
[  2.00029717e-01,  4.64788881e-01,  1.21890937e+00],
[ -5.15103652e-01, -1.09058306e+00, -3.45582490e-01],
[  1.45715009e+00, -1.18475597e-01,  2.91803083e-01],
[  8.35145683e-01,  1.33968560e+00,  5.81523798e-01],
[  1.01321598e+00, -1.18779381e+00, -1.15680049e+00],
[ -1.09283526e+00,  1.04805336e+00,  4.65635512e-01],
[  1.19747311e+00,  1.82573933e+00,  1.50863009e+00],
[ -9.70148534e-01, -4.10107836e-01, -1.30166085e+00],
[ -1.76683183e+00, -3.12897090e-01, -3.74554562e-01],
[ -1.70662366e+00, -4.10107836e-01,  1.30582558e+00],
[ -1.15694568e+00,  2.02016082e+00,  5.23579655e-01],
[  7.23725813e-01,  6.59210374e-01, -1.09885635e+00],
[  7.62251660e-01, -8.96161568e-01,  3.78719297e-01],
[ -1.70893663e+00, -1.18779381e+00,  2.91803083e-01],
[  1.44193841e+00,  1.04805336e+00, -1.21474464e+00],
[ -7.36808194e-01, -1.47942605e+00, -1.44652121e+00],
[ -1.39522414e+00, -6.04529329e-01, -1.50446535e+00],
[  1.42143386e+00,  2.11737157e+00, -8.09135634e-01],
[  1.60063936e+00, -1.86826903e+00,  1.75914797e-01],
[  1.19444714e-01, -2.15686344e-01,  8.42272441e-01],
[ -4.13418121e-01, -1.86826903e+00, -1.27268878e+00],
[  1.18546837e+00,  2.11737157e+00,  3.78719297e-01],
[  7.83579228e-01, -1.38221530e+00,  5.52551726e-01],
[  5.33139939e-01, -1.09058306e+00, -3.45582490e-01],
[ -1.21855865e+00,  1.73156642e-01, -6.64275277e-01],
[  3.78752752e-01,  3.67578135e-01,  2.08236764e-03],
[  5.14746197e-02, -6.04529329e-01,  2.31984809e+00],
[  1.70245260e+00, -3.12897090e-01,  2.04886868e-01],
[ -1.61799293e+00, -1.57663679e+00, -2.00722133e-01],
[ -5.91928306e-01,  6.59210374e-01, -1.38857706e+00],
[  1.40000697e+00, -1.09058306e+00,  5.52551726e-01],
[ -8.83036131e-01, -1.96547978e+00,  3.49747226e-01],
[  1.09392869e+00,  3.67578135e-01,  2.62831011e-01],
[ -1.32153547e+00,  1.73156642e-01, -2.87638347e-01],
[  1.24344871e+00,  1.43689635e+00, -1.04091221e+00],
[ -7.32565384e-01,  8.53631867e-01,  1.07404901e+00]]])
```

```
In [9]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(xtrain, ytrain)
```

```
/home/rmdstic/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)
```

```
Out[9]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
        intercept_scaling=1, max_iter=100, multi_class='warn',
        n_jobs=None, penalty='l2', random_state=None, solver='warn',
        tol=0.0001, verbose=0, warm_start=False)
```

```
In [10]: ytrain_pred = lr.predict(xtrain)
ytest_pred = lr.predict(xtest)
```

In [11]: ytrain_pred

```
Out[11]: array([[1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1,
0, 0,
0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1,
1, 1,
0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0, 1,
0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0,
0, 0,
0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0,
1, 1,
0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0,
0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1,
0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0,
1, 0,
1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0,
0, 0,
1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0,
0, 0,
0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0,
1, 0,
0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0,
1, 0,
0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0,
0, 1,
1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]])
```

In [12]: ytest_pred

```
Out[12]: array([[0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0, 1,
0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0,
1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 1,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1])
```

In [13]: print(lr.predict([[15624510,19,19000]]))
[0]

In [14]: lr.predict([[4.88966473e-01,-7.98950822e-01,4.94607583e-01]])

Out[14]: array([0])

In [15]: lr.predict([[-1.50619343e-01,-2.15686344e-01,2.14601566e+00]])

Out[15]: array([1])

```
In [17]: from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
matrix = confusion_matrix(ytest, ytest_pred)
print(matrix)
```

```
[[56  2]
 [ 5 17]]
```

```
In [18]: score = accuracy_score(ytest, ytest_pred)
score
```

Out[18]: 0.9125

```
In [19]: cr = classification_report(ytest, ytest_pred)
print(cr)
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.97 | 0.94 | 58 |
| 1 | 0.89 | 0.77 | 0.83 | 22 |
| micro avg | 0.91 | 0.91 | 0.91 | 80 |
| macro avg | 0.91 | 0.87 | 0.89 | 80 |
| weighted avg | 0.91 | 0.91 | 0.91 | 80 |

In [20]: `pip install seaborn`

```
Requirement already satisfied: seaborn in /home/rmdstic/anaconda3/lib/python3.7/site-packages (0.9.0)
Requirement already satisfied: matplotlib>=1.4.3 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from seaborn) (3.0.3)
Requirement already satisfied: pandas>=0.15.2 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from seaborn) (0.24.2)
Requirement already satisfied: numpy>=1.9.3 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from seaborn) (1.16.2)
Requirement already satisfied: scipy>=0.14.0 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from seaborn) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from matplotlib>=1.4.3->seaborn) (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from matplotlib>=1.4.3->seaborn) (1.0.1)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from matplotlib>=1.4.3->seaborn) (2.3.1)
Requirement already satisfied: python-dateutil>=2.1 in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from matplotlib>=1.4.3->seaborn) (2.8.0)
Requirement already satisfied: pytz>=2011k in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from pandas>=0.15.2->seaborn) (2018.9)
Requirement already satisfied: six in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from cycler>=0.10->matplotlib>=1.4.3->seaborn) (1.12.0)
Requirement already satisfied: setuptools in /home/rmdstic/anaconda3/lib/python3.7/site-packages (from kiwisolver>=1.0.1->matplotlib>=1.4.3->seaborn) (40.8.0)
Note: you may need to restart the kernel to use updated packages.
```

In [21]: `import seaborn as sns`
`sns.heatmap(matrix,annot=True)`

Out[21]: `<matplotlib.axes._subplots.AxesSubplot at 0x7f0e805f3ba8>`



In [22]: `tn, fp, fn, tp = confusion_matrix(ytest,ytest_pred).ravel()`

In [23]: `print(tn, fp, fn, tp)`

56 2 5 17

In [24]: `print("Error rate:", (fp+fn)/(tn+fp+fn+tp))`

Error rate: 0.0875

In []: